
CS267 Lecture 2

Single Processor Machines: Memory Hierarchies and Processor Features

Case Study: Tuning Matrix Multiply

Kathy Yelick

<https://sites.google.com/lbl.gov/cs267-spr2018/>

Motivation

- Parallel computing is about performance
 - Distributed computing is a different thing
- Most applications run at $< 10\%$ of the “peak” performance
 - Peak is the maximum the hardware can physically execute
- Much of the performance is lost on a single processor
 - The code running on one processor often runs at only 10-20% of the processor peak
- Most of that loss is in the memory system
 - Moving data takes much longer than arithmetic and logic
- We need to look under the hood of modern processors
 - For today, we will look at only a single “core” processor
 - These issues will exist on processors within any parallel computer

Possible conclusions to draw from today's lecture


- “Computer architectures are fascinating, and I really want to understand why apparently simple programs can behave in such complex ways!”
- “I want to learn how to design algorithms that run really fast no matter how complicated the underlying computer architecture.”
- “I hope that most of the time I can libraries or other software that hides all these details from me so I don't have to worry about them!”
- “I wish the compiler would handle everything for me.”
- “I would like to write a compiler that would handle all of these details.”
- I want to understand how Meltdown/Spectre work

Outline

- Costs in modern processors
 - Idealized models and actual costs
- Memory hierarchies
- Parallelism in a single processor
- Case study: Matrix multiplication

Idealized Uniprocessor Model


- Processor names **variables**:
 - Integers, floats, pointers, arrays, structures, etc.
- Processor performs **operations** on those variables:
 - Arithmetic, logical operations, etc.
- Processor **controls** the order, as specified by program
 - Branches (if), loops, function calls, etc.



- **Idealized Cost**
 - Each operation has roughly the same cost
add, multiply, etc.


Slightly Less Idealized Uniprocessor Model

- Processor names **variables**:
 - Integers, floats, pointers, arrays, structures, etc.
 - These are really words, e.g., 64-bit doubles, 32-bit ints, bytes, etc.
- Processor performs **operations** on those variables:
 - Arithmetic, logical operations, etc.
 - Only performs these operations on values in registers
- Processor **controls** the order, as specified by program
 - Branches (if), loops, function calls, etc.



- ***Idealized Cost***
 - Each operation has roughly the same cost
add, multiply, etc.



Compilers and assembly code



- Compilers for languages like C/C++ and Fortran:
 - Check that the program is legal
 - Translate into assembly code
 - Optimizes the generated code

Compilers Manage Memory and Registers

- Compiler performs “register allocation” to decide when to load/store and when to reuse



Register allocation in first Fortran compiler in 1950s, graph coloring in 1980. JITs may use something cheaper

Compiler Optimizes Code

- **Compiler performs a number of optimizations:**
 - Unrolls loops (because control isn't free)
 - Fuses loops (merge two together)
 - Interchanges loops (reorder)
 - Eliminates dead code (the branch never taken)
 - Reorders instructions to improve register reuse and more
 - Strength reduction (turns expensive instruction, e.g., multiply by 2, into cheaper one, shift left)
- **The compiler managed operations (load/store) that are important to performance but we don't "see" in code.**
- **Why is this your problem?**
 - Because sometimes it does the best thing possible
 - But other times it does not...

Outline


- Idealized and actual costs in modern processors
- Memory hierarchies
 - Temporal and spatial locality
 - Basics of caches
 - Use of microbenchmarks to characterized performance
- Parallelism within single processors
- Case study: Matrix Multiplication

Outline

- Costs in modern processors
- Memory hierarchies
 - Temporal and spatial locality
 - Basics of caches
 - Use of microbenchmarks to characterize performance
- Parallelism in a single processor
- Case study: Matrix multiplication

Memory Bandwidth Gap

Memory Bandwidth is Falling Behind: $(\text{GFLOP/s}) / (\text{GWord/s})$



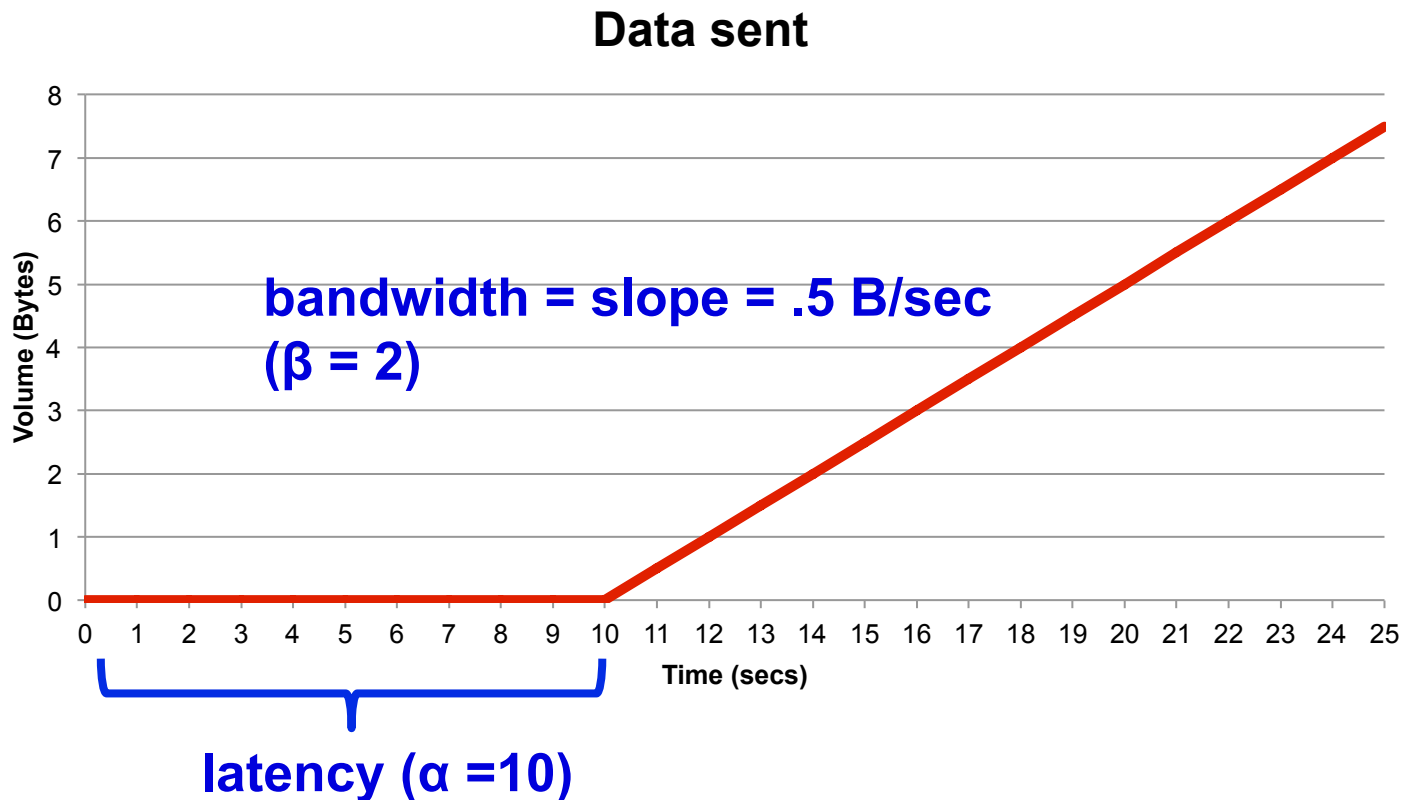
+14.2%/year
2x / 5.2 years

RISC systems (IBM, MIPS, Alpha)

x86-64 systems (AMD & Intel)

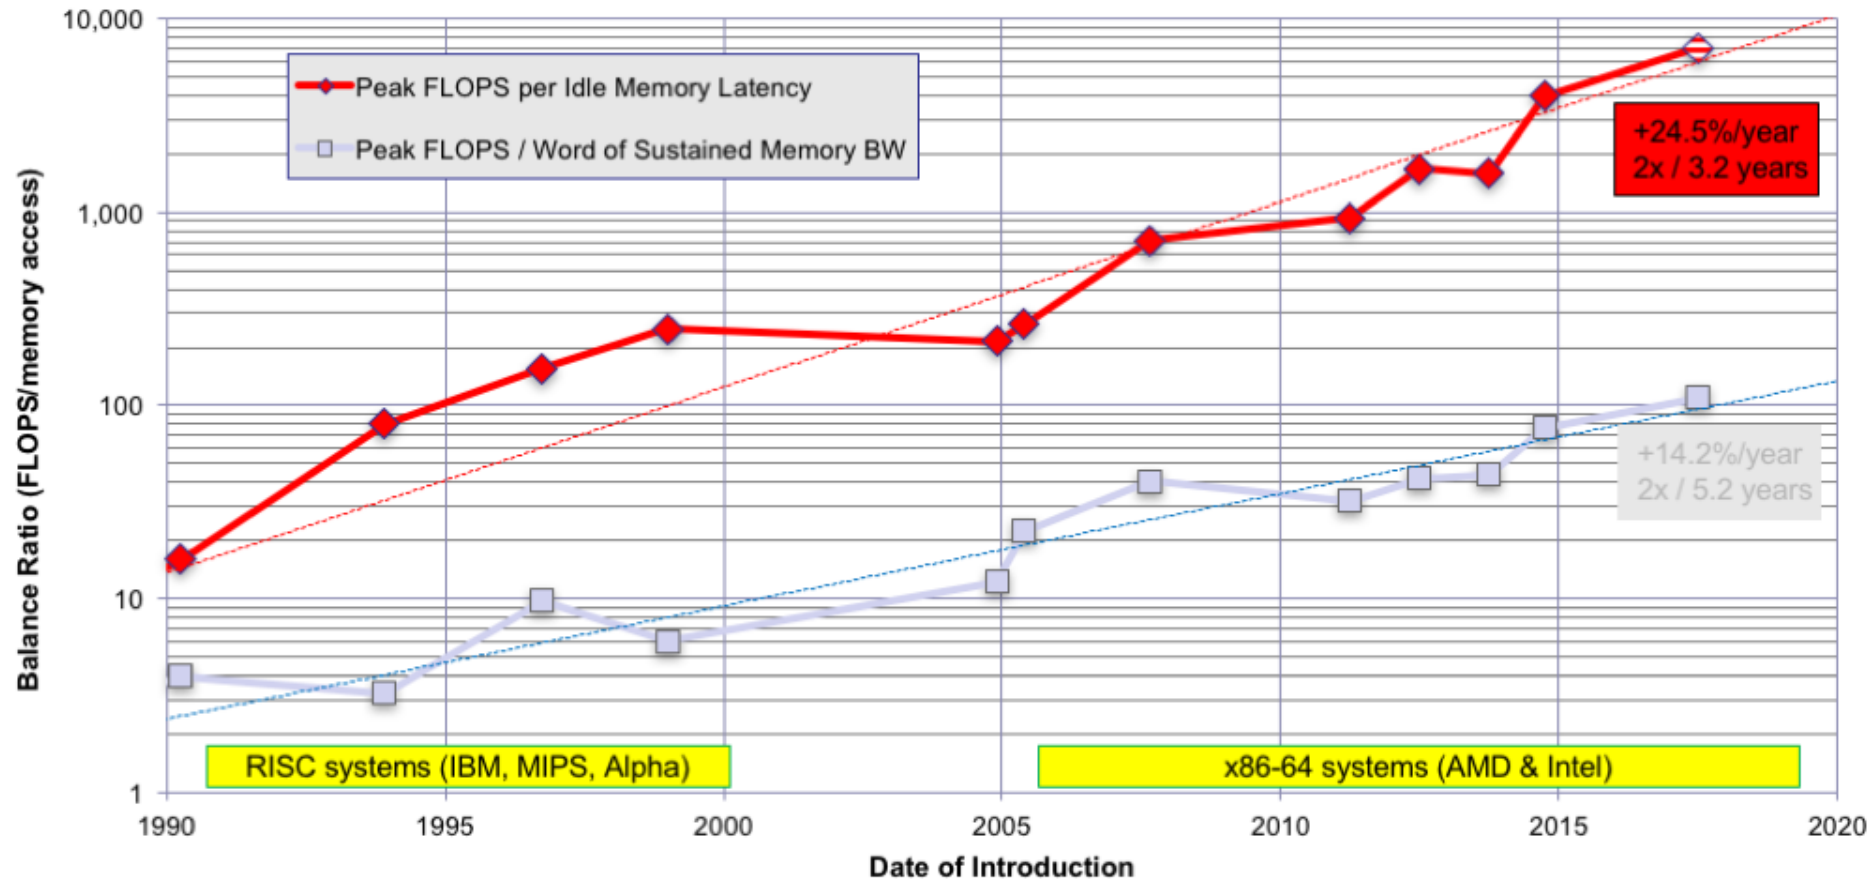
Latency vs Bandwidth

- Previous graph show bandwidth gap
- Many applications are limited not by memory bandwidth, but by latency



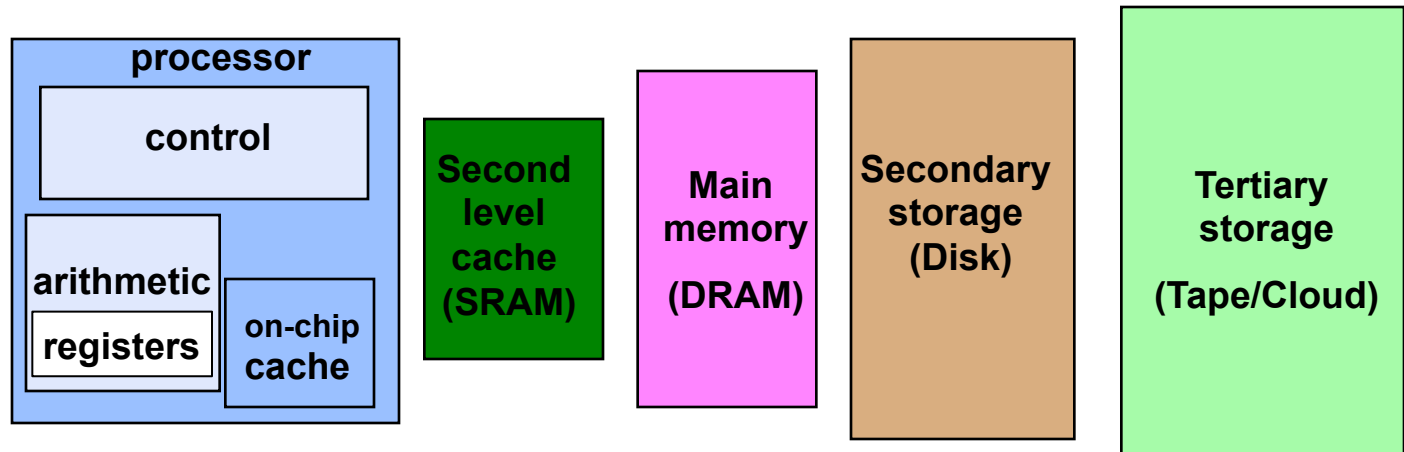
Memory Latency Gap is Worse

Memory Latency is much worse: $(\text{GFLOP/s}) / (\text{Memory Latency})$



Memory Hierarchy

- Most programs have a high degree of locality
 - **spatial locality**: accessing things nearby previous accesses
 - **temporal locality**: reusing an item that was previously accessed
- Memory hierarchy use this to improve *average case*



Speed	1ns	10ns	100ns	10ms	10sec
Size	KB	MB	GB	TB	PB

Cache Basics

- **Cache** is fast (expensive) memory which keeps copy of data in main memory; it is hidden from software
 - **Simplest example:** data at memory address **xxxxx1101** is stored at cache location **1101**
- **Cache hit:** in-cache memory access—cheap
- **Cache miss:** non-cached memory access—expensive
 - **Need to access next, slower level of cache**
- **Cache line length:** # of bytes loaded together in one entry
 - **Ex:** If either **xxxxx1100** or **xxxxx1101** is loaded, both are
- **Associativity**
 - **direct-mapped:** only 1 address (line) in a given range in cache
 - Data stored at address **xxxxx1101** stored at cache location **1101**, in 16 word cache
 - **n -way:** $n \geq 2$ lines with different addresses can be stored
 - Up to $n \leq 16$ words with addresses **xxxxx1101** can be stored at cache location **1101** (so cache can store $16n$ words)

Why Have Multiple Levels of Cache?

- On-chip vs. off-chip
 - **On-chip caches are faster, but limited in size**
- A large cache has delays
 - **Hardware to check longer addresses in cache takes more time**
 - **Associativity, which gives a more general set of data in cache, also takes more time**
- Some examples:
 - **Cray T3E eliminated one cache to speed up misses**
 - **Intel Haswell (Cori Ph1) uses a Level 4 cache as “victim cache”**
- There are other levels of the memory hierarchy
 - **Register, pages (TLB, virtual memory), ...**
 - **And it isn't always a hierarchy**

Approaches to Handling Memory Latency

- Eliminate memory operations by saving values in small, fast memory (cache or registers) and reusing them (bandwidth filtering)
 - **need temporal locality in program**
- Take advantage of better bandwidth by getting a chunk of memory into cache (or vector registers) and using whole chunk
 - **need spatial locality in program**
- Take advantage of better bandwidth by allowing processor to issue multiple reads or writes with a single instruction
 - **vector operations require access set of locations (typically neighboring), requires they are independent**
- Take advantage of better bandwidth by allowing processor to issue reads/writes in parallel with other reads/writes/operations
 - **prefetching issues read hint**
 - **delayed writes (write buffering) stages writes for later operation**
 - **both require that nothing dependent is happening (parallelism)**

concurrency

How much concurrency do you need?

- To run at bandwidth speeds rather than latency
- **Little's Law** from queuing theory says:

$$\text{concurrency} = \text{latency} * \text{bandwidth}$$

- In my earlier example with


latency = 10 sec

bandwidth = 2 Bytes/sec

- Requires 20 bytes in flight concurrently to reach bandwidth speeds
- That means finding 20 independent things to issue

More realistic example

Little's Law: illustration for 2005-era Opteron processor
60 ns latency, 6.4 GB/s (=10ns per 64B cache line)




- $60 \text{ ns} * 6.4 \text{ GB/s} = 384 \text{ Bytes} = 6 \text{ cache lines}$
- To keep the pipeline full, there must always be 6 cache lines “in flight”
- Each request must be launched at least 60 ns before the data is needed

Stream benchmark for measuring bandwidth

- Stream benchmark (also due to McCalpin)
- Four kernels, all for $i = 1$ to N :
 - **Copy:** $C[i] = A[i];$
 - **Scale:** $B[i] = \text{scalar} * C[i];$
 - **Add:** $C[i] = A[i] + B[i];$
 - **Triad:** $A[i] = B[i] + \text{scalar} * C[i];$
- N chose so that array is much larger than cache
- Repeated ~ 10 times (ignoring first) to “warm cache
- Min/Av/Max timing report

Experimental Study of Memory (Membench)


- Microbenchmark for memory system performance



- for array A of length L from 4KB to 8MB by 2x
for stride s from 4 Bytes (1 word) to L/2 by 2x
time the following loop
(repeat many times and average)
for i from 0 to L-1 by s
load A[i] from memory (4 Bytes)


1 experiment

Membench: What to Expect



- Consider the average cost per load
 - Plot one line for each array length, time vs. stride
 - Small stride is best: if cache line holds 4 words, at most $\frac{1}{4}$ miss
 - If array is smaller than a given cache, all those accesses will hit (after the first run, which is negligible for large enough runs)
 - Picture assumes only one level of cache
 - Values have gotten more difficult to measure on modern procs

Memory Hierarchy on a Sun Ultra-2i



See www.cs.berkeley.edu/~yelick/arvindk/t3d-isca95.ps for details

Memory Hierarchy on a Power3 (Seaborg)


Power3, 375 MHz

Saavedra-Barrera Benchmark: Time to execute 1 load [Power3]

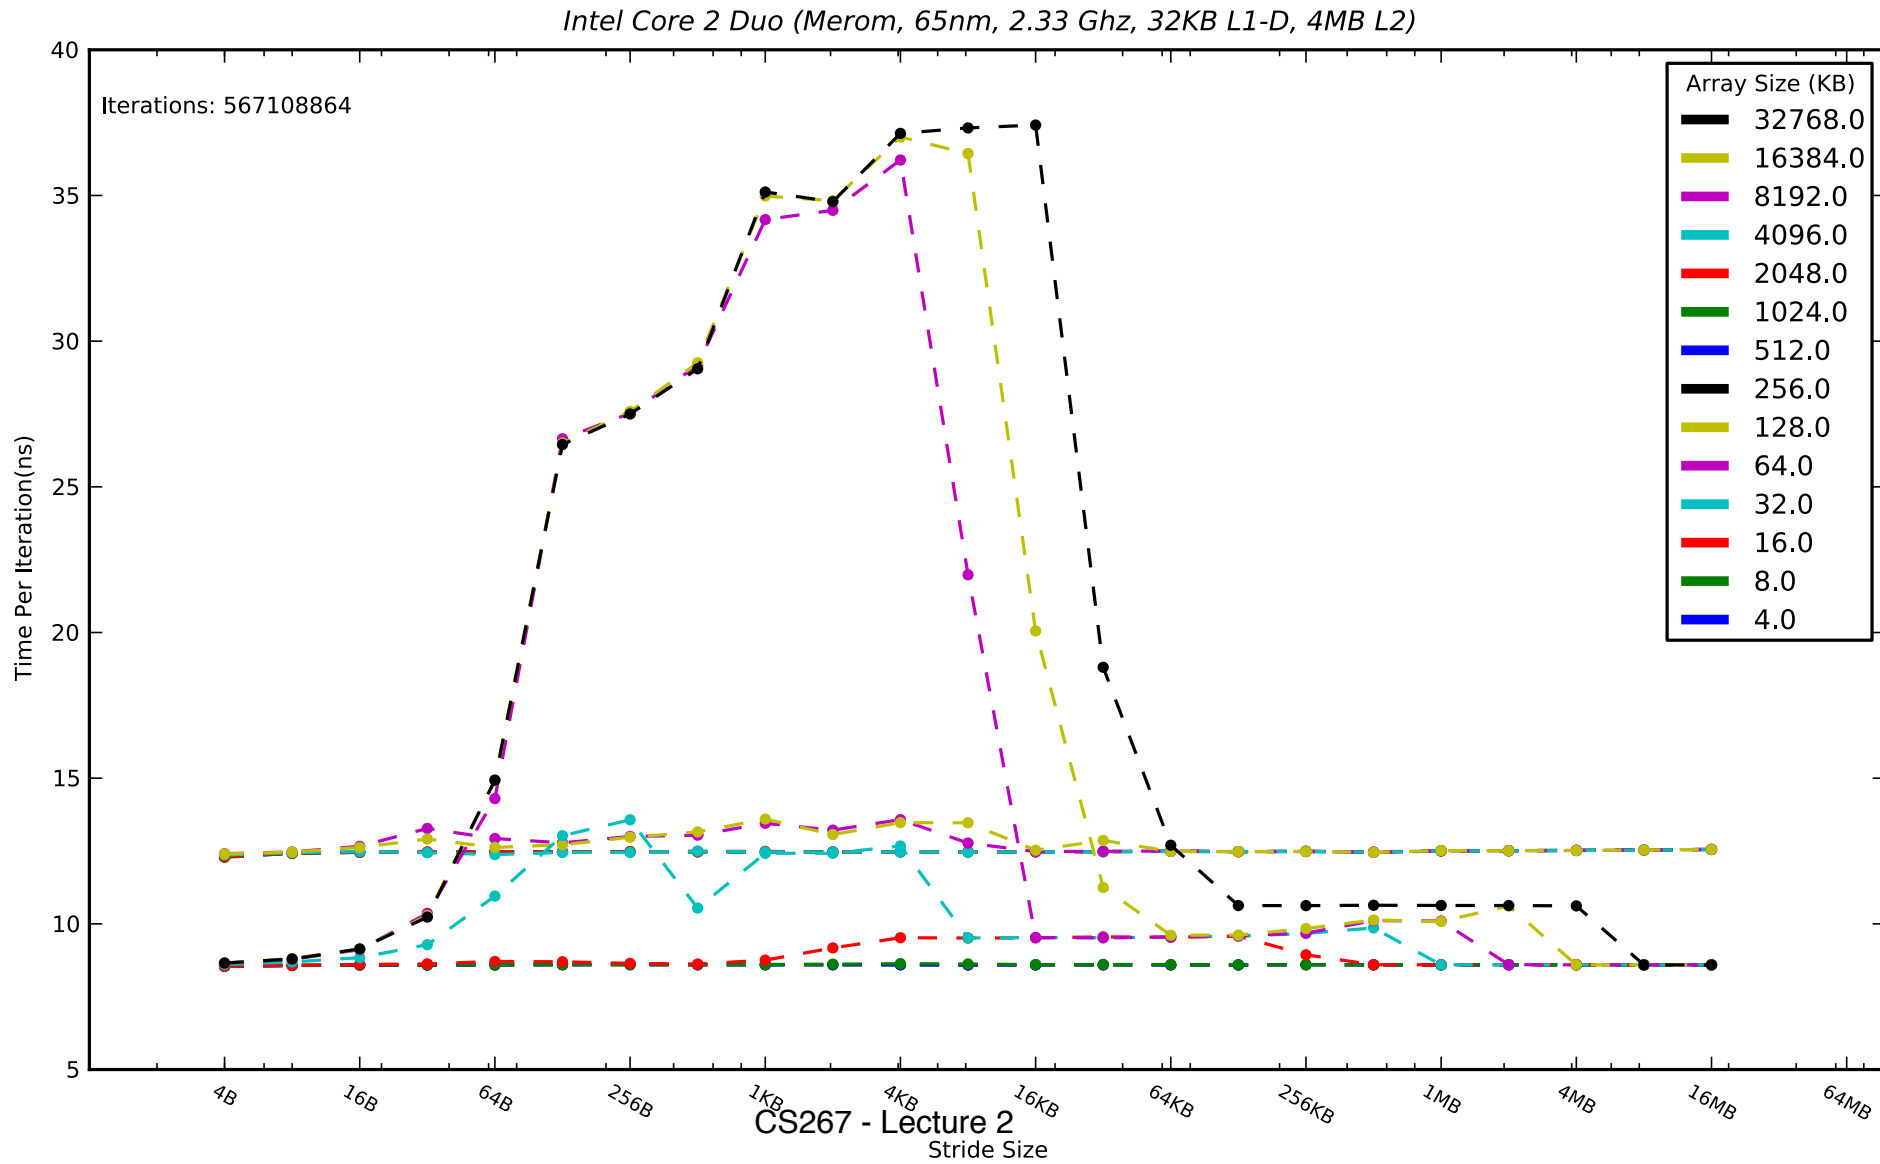
Mem: 396 ns
(132 cycles)

L2: 8 MB
128 B line
9 cycles

L1: 32 KB
128B line
.5-2 cycles




Memory Hierarchy on an Intel Core 2 Duo




Stanza Triad

- Even smaller benchmark for prefetching
- Derived from STREAM Triad
- Stanza (L) is the length of a unit stride run

```
while i < arraylength
  for each  $L$  element stanza
     $A[i] = \text{scalar} * X[i] + Y[i]$ 
  skip  $k$  elements
```

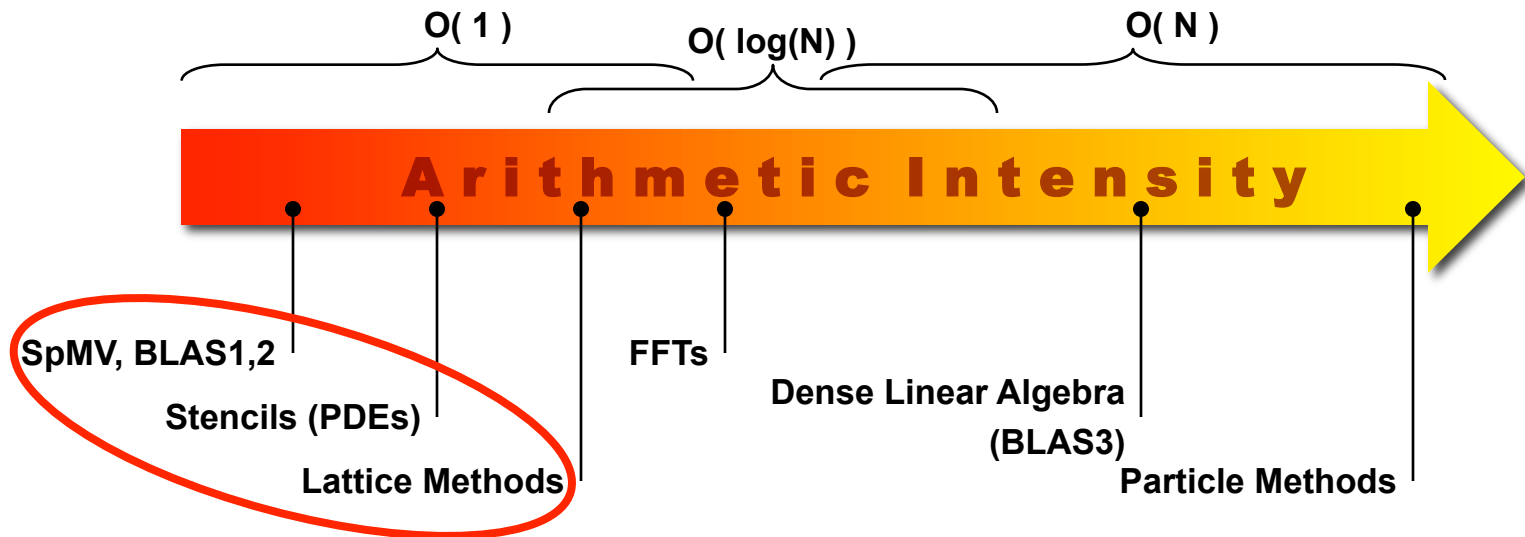


Stanza Triad Results – Prefetch matters



- This graph (x-axis) starts at a cache line size (≥ 16 Bytes)
- If cache locality was the only thing that mattered, we would expect
 - Flat lines equal to measured memory peak bandwidth (STREAM) as on Pentium3
- Prefetching gets the next cache line (pipelining) while using the current one
 - This does not “kick in” immediately, so performance depends on L
 - http://crd-legacy.lbl.gov/~oliker/papers/msp_2005.pdf


Arithmetic Intensity



- Arithmetic Intensity (AI) \sim Total Flops / Total DRAM Bytes moved
- Can we hit the bandwidth or flop limit?
- The rest is all about having and being able to utilize enough concurrency

The Roofline Performance Model

Log scale!




flop:DRAM byte ratio

Log scale

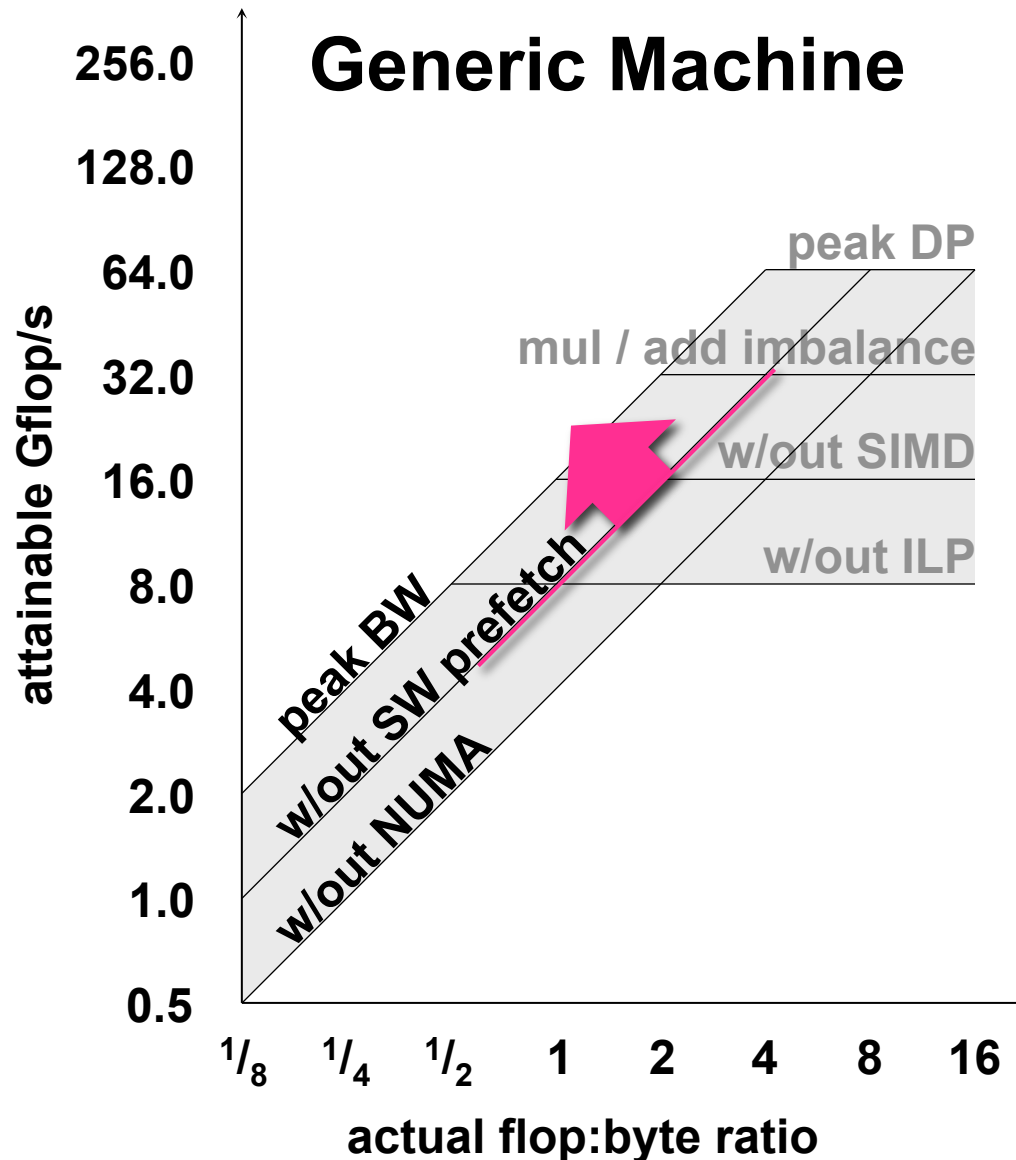
Sam Williams PhD, 2008

The Roofline Performance Model



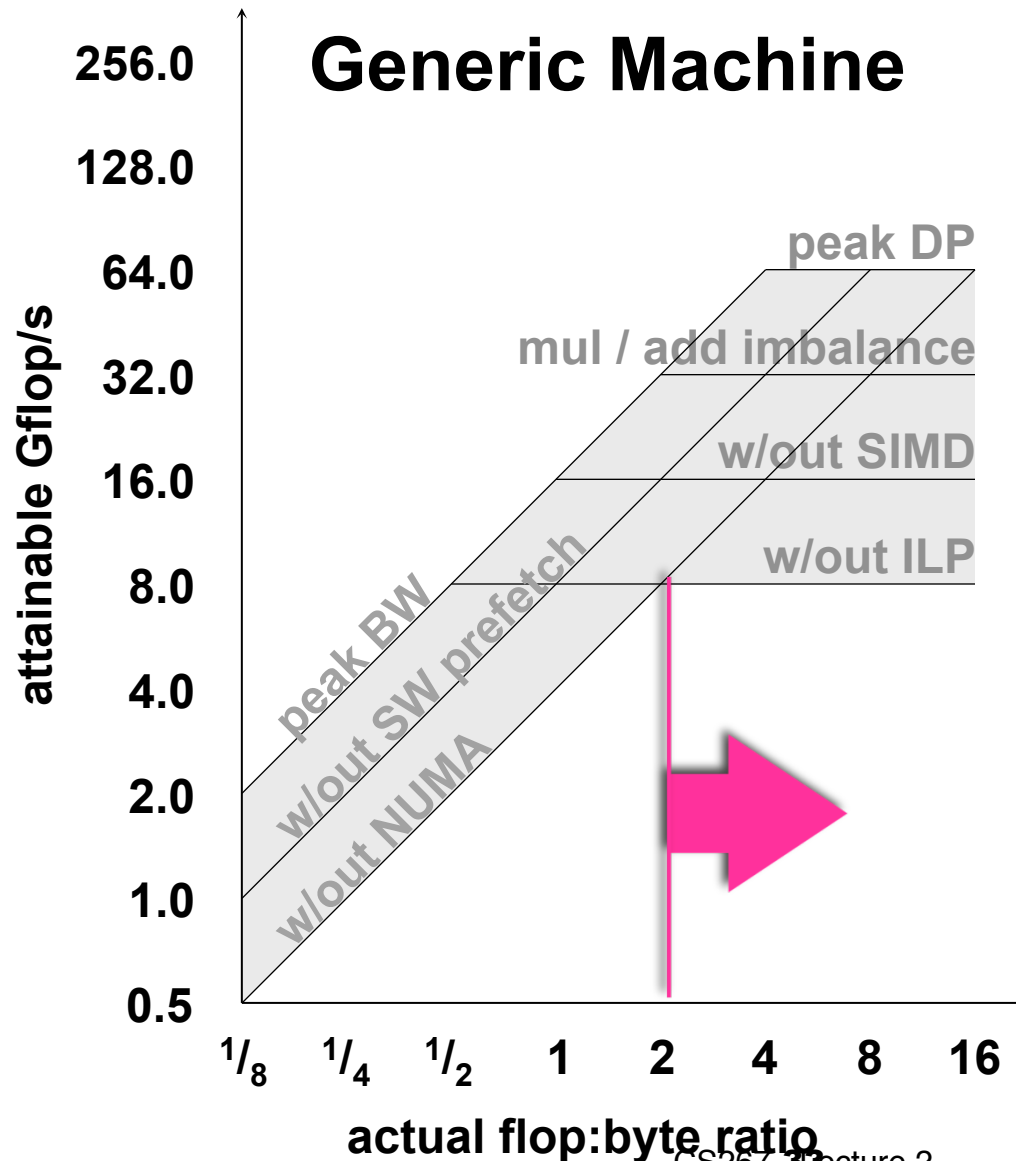
- ❖ The top of the roof is determined by peak computation rate (Double Precision floating point, DP for these algorithms)
- ❖ The instruction mix, lack of SIMD operations, ILP or failure to use other features of peak will lower attainable

The Roofline Performance Model



- ❖ The sloped part of the roof is determined by peak DRAM bandwidth (STREAM)
- ❖ Lack of proper prefetch, ignoring NUMA, or other things will reduce attainable bandwidth

The Roofline Performance Model



- ❖ Locations of posts in the building are determined by algorithmic intensity
- ❖ Will vary across algorithms and with bandwidth-reducing optimizations, such as better cache re-use (tiling), compression techniques

Lessons

- Actual performance of a simple program can be a complicated function of the architecture
 - Slight changes in the architecture or program change the performance significantly
 - To write fast programs, need to consider architecture
 - True on sequential or parallel processor
 - We would like simple models to help us design efficient algorithms
- We will illustrate with a common technique for improving cache performance, called **blocking** or **tiling**
 - Idea: used divide-and-conquer to define a problem that fits in register/L1-cache/L2-cache

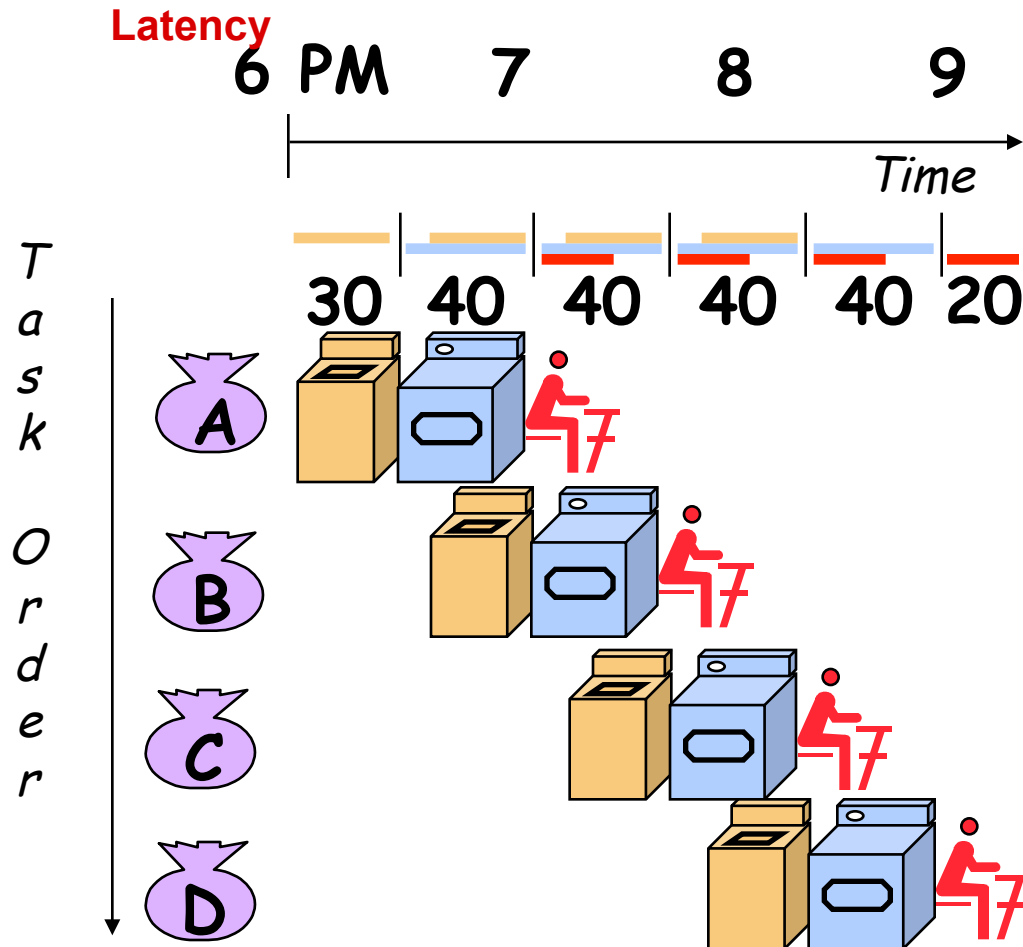
Outline

- Idealized and actual costs in modern processors
- Memory hierarchies
- Parallelism within single processors
 - Hidden from software (sort of)
 - Pipelining
 - SIMD units
- Case study: Matrix Multiplication

What is Pipelining?

Dave Patterson's Laundry example: 4 people doing laundry

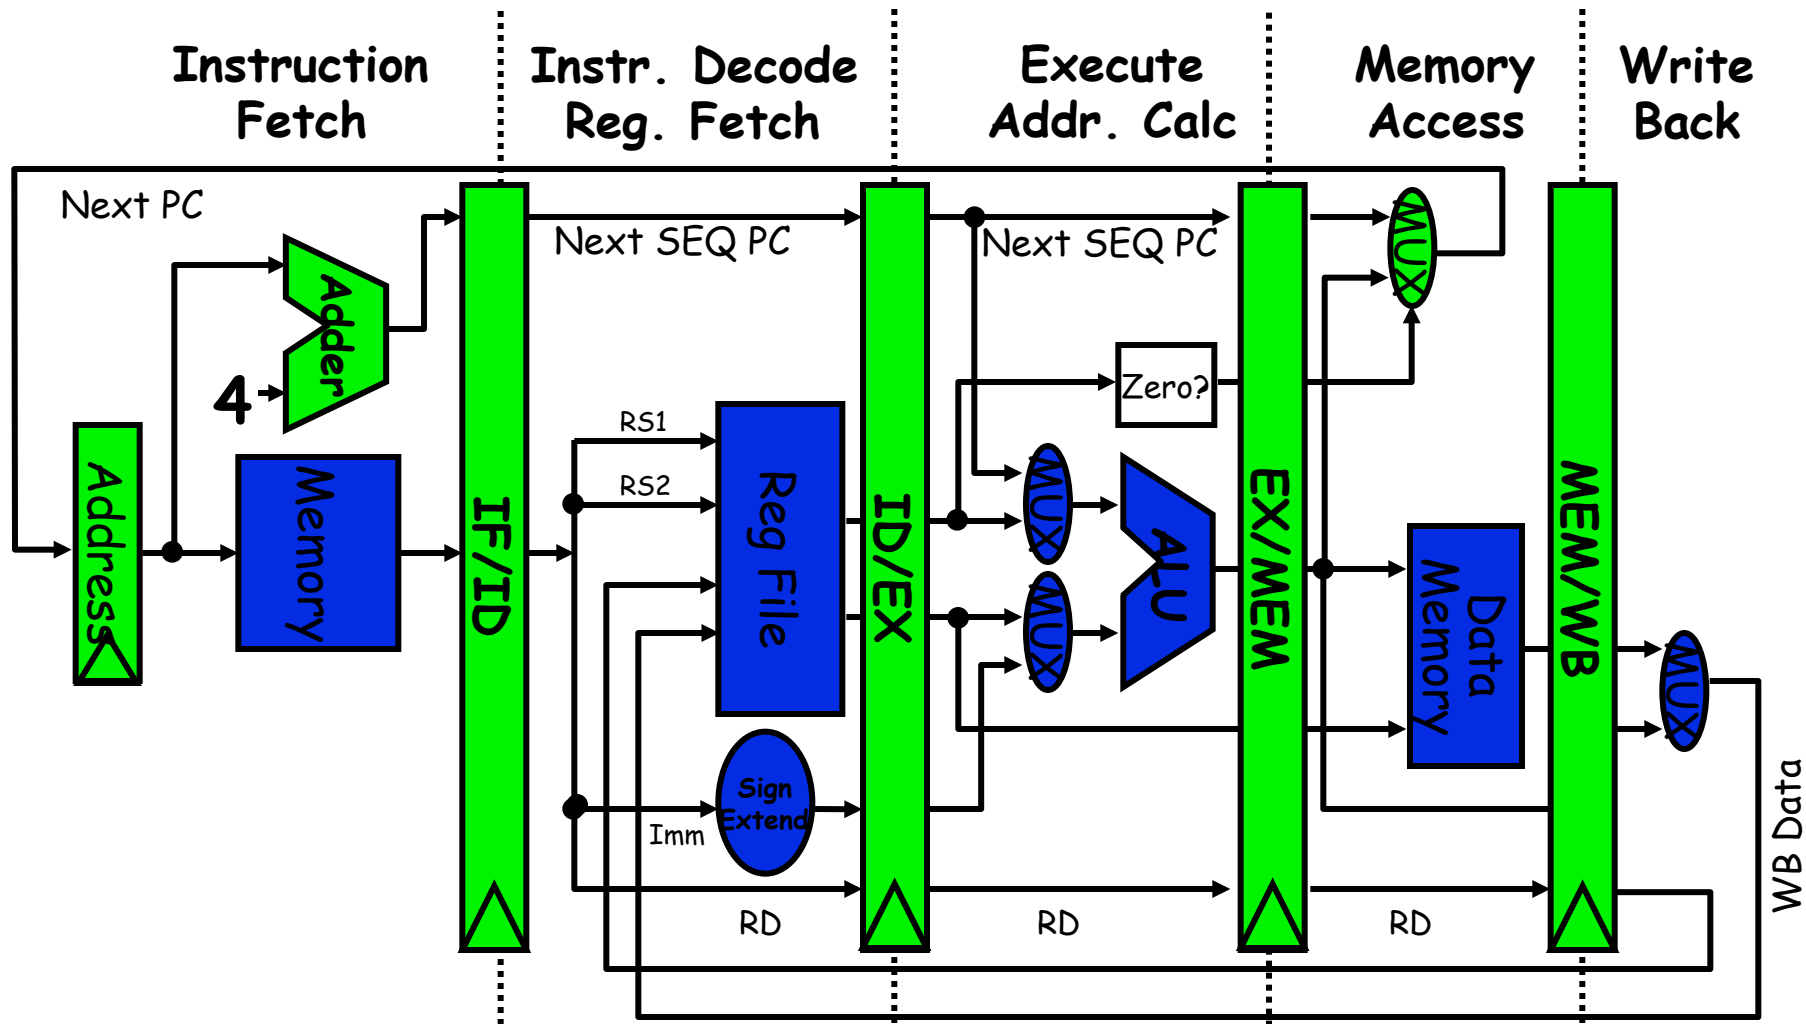
wash (30 min) + dry (40 min) + fold (20 min) = 90 min



- In this example:
 - Sequential execution takes $4 * 90\text{min} = 6 \text{ hours}$
 - Pipelined execution takes $30 + 4 * 40 + 20 = 3.5 \text{ hours}$
- **Bandwidth** = loads/hour
- $BW = 4/6 \text{ l/h}$ w/o pipelining
- $BW = 4/3.5 \text{ l/h}$ w pipelining
- $BW \leq 1.5 \text{ l/h}$ w pipelining, more total loads
- Pipelining doesn't change **latency** (90 min)
- Bandwidth limited by **slowest** pipeline stage
- Speedup \leq # of stages

Example: 5 Steps of MIPS Datapath

Figure 3.4, Page 134 , CA:AQA 2e by Patterson and Hennessy



- Pipelining is also used within arithmetic units
 - a fp multiply may have latency 10 cycles, but throughput of 1/cycle

SIMD: Single Instruction, Multiple Data

Scalar processing

- traditional mode
- one operation produces one result



+



SIMD processing: vectors

- Sandy Bridge: AVX (256 bit)
- Haswell: AVX2 (256 bit w/ FMA)
- KNL: AVX-512 (512 bit)

X



+

Y




X + Y



SSE / SSE2 SIMD on Intel

- SSE2 data types: anything that fits into 16 bytes, e.g.,



- Instructions perform add, multiply etc. on all the data in this 16-byte register in parallel
- Challenges:
 - Need to be contiguous in memory and aligned
 - Some instructions to move data around from one part of register to another
- Similar on GPUs, vector processors (but many more simultaneous operations)

What does this mean to you?

- In addition to SIMD extensions, the processor may have other special instructions
 - Fused Multiply-Add (FMA) instructions:
$$x = y + c * z$$
is so common some processor execute the multiply/add as a single instruction, at the same rate (bandwidth) as + or * alone
- In theory, the compiler understands all of this
 - When compiling, it will rearrange instructions to get a good “schedule” that maximizes pipelining, uses FMAs and SIMD
 - It works with the mix of instructions inside an inner loop or other block of code
- But in practice the compiler may need your help
 - Choose a different compiler, optimization flags, etc.
 - Rearrange your code to make things more obvious
 - Using special functions (“intrinsics”) or write in assembly ☹

Outline

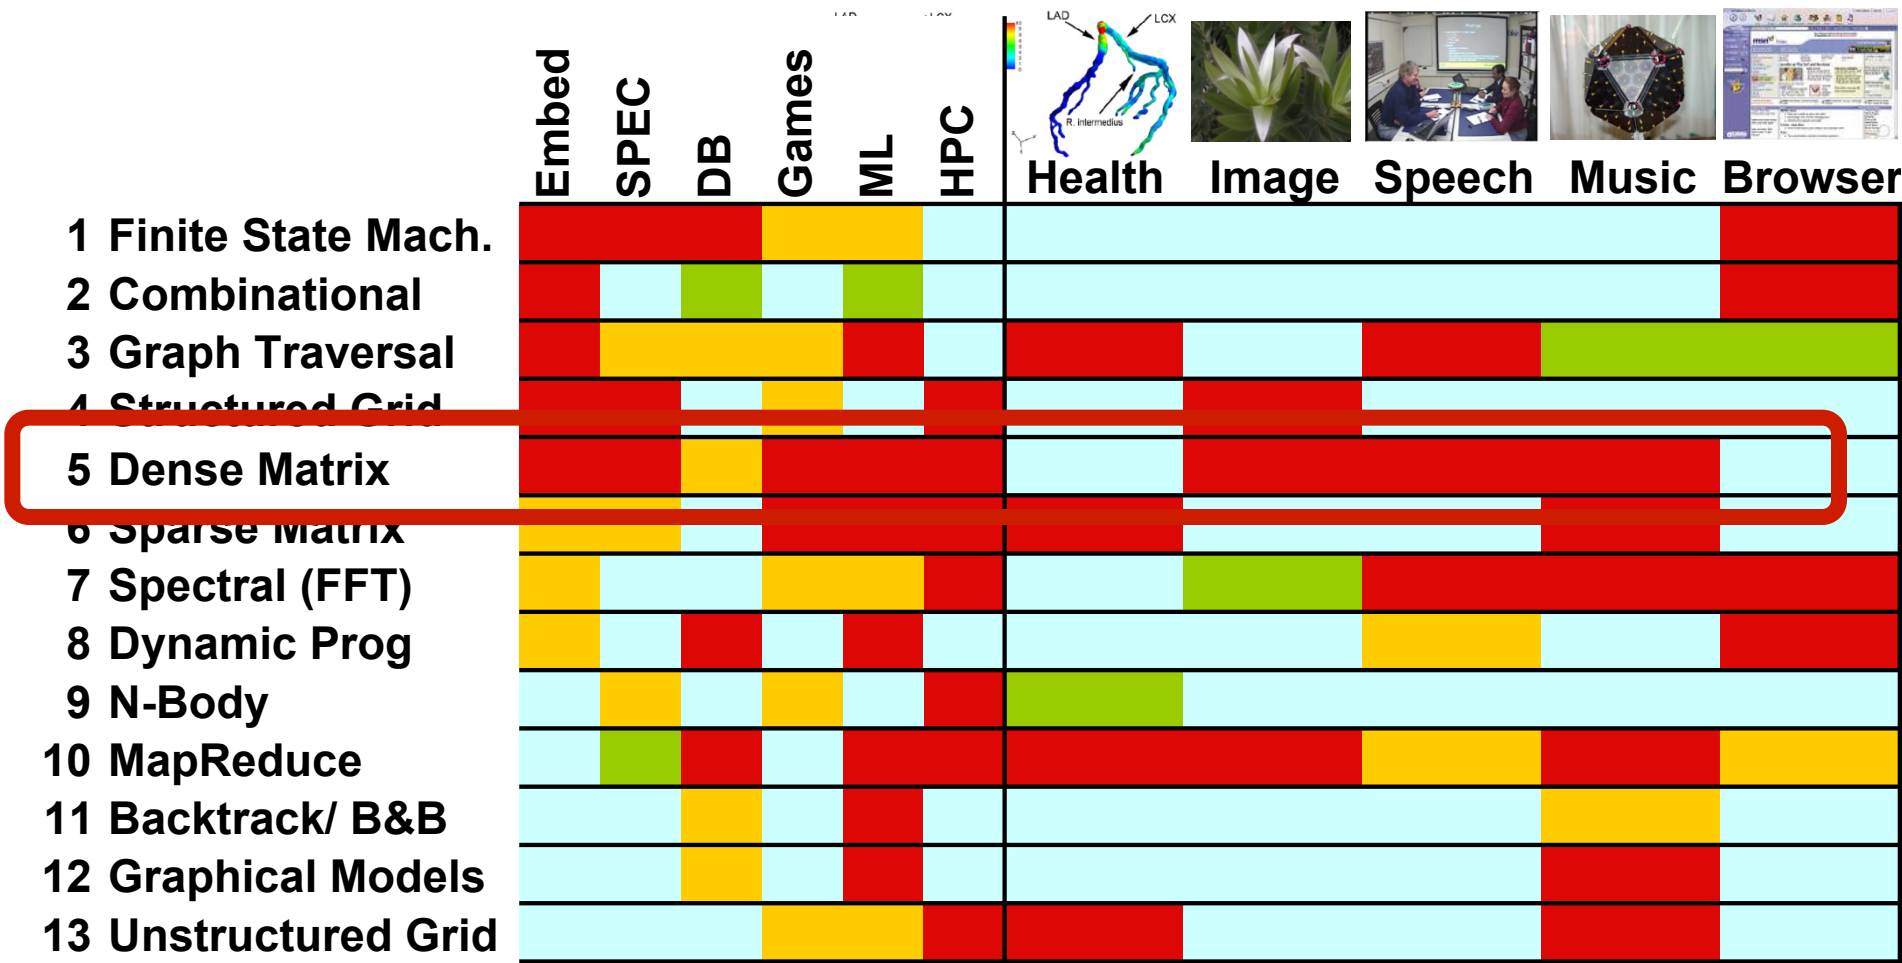
- Idealized and actual costs in modern processors
- Memory hierarchies
- Parallelism within single processors
- Case study: Matrix Multiplication
 - Use of performance models to understand performance
 - Attainable lower bounds on communication
 - Simple cache model
 - Warm-up: Matrix-vector multiplication
 - Naïve vs optimized Matrix-Matrix Multiply
 - Minimizing data movement
 - Beating $O(n^3)$ operations
 - Practical optimizations (*continued next time*)

Why Matrix Multiplication?


- An important kernel in many problems
 - Appears in many linear algebra algorithms
 - Bottleneck for dense linear algebra, including Top500
 - One of the motifs of parallel computing
 - Closely related to other algorithms, e.g., transitive closure on a graph using Floyd-Warshall
- Optimization ideas can be used in other problems
- The best case for optimization payoffs
- The most-studied algorithm in high performance computing

What do commercial and CSE applications have in common?

Motif/Dwarf: Common Computational Methods (Red Hot → Blue Cool)



Matrix-multiply, optimized several ways



Speed of n-by-n matrix multiply on Sun Ultra-1/170, peak = 330 MFlops

Note on Matrix Storage

- A matrix is a 2-D array of elements, but memory addresses are “1-D”
- Conventions for matrix layout
 - by column, or “column major” (Fortran default); $A(i,j)$ at $A+i*j*n$
 - by row, or “row major” (C default) $A(i,j)$ at $A+i*n+j$
 - recursive (later)

Column major

↓


0	5	10	15
1	6	11	16
2	7	12	17
3	8	13	18
4	9	14	19

Row major

→

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
16	17	18	19

Column major matrix in memory



- Column major (for now)

Using a Simple Model of Memory to Optimize

- Assume just 2 levels in the hierarchy, fast and slow
- All data initially in slow memory
 - m = number of memory elements (words) moved between fast and slow memory
 - t_m = time per slow memory operation
 - f = number of arithmetic operations
 - t_f = time per arithmetic operation $\ll t_m$
 - $q = f / m$ average number of flops per slow memory access
- Minimum possible time = $f * t_f$ when all data in fast memory
- Actual time
 - $f * t_f + m * t_m = f * t_f * (1 + \boxed{t_m/t_f} * 1/q)$
- Larger q means time closer to minimum $f * t_f$
 - $q \geq t_m/t_f$ needed to get at least half of peak speed

Computational Intensity: Key to algorithm efficiency

Machine Balance: Key to machine efficiency


Warm up: Matrix-vector multiplication

{implements $y = y + A*x$ }

for $i = 1:n$

 for $j = 1:n$

$y(i) = y(i) + A(i,j)*x(j)$



Warm up: Matrix-vector multiplication

```
{read x(1:n) into fast memory}  
{read y(1:n) into fast memory}  
for i = 1:n  
    {read row i of A into fast memory}  
    for j = 1:n  
        y(i) = y(i) + A(i,j)*x(j)  
    {write y(1:n) back to slow memory}
```

- m = number of slow memory refs = $3n + n^2$
- f = number of arithmetic operations = $2n^2$
- $q = f / m \approx 2$
- Matrix-vector multiplication limited by slow memory speed

Modeling Matrix-Vector Multiplication

- Compute time for $n \times n = 1000 \times 1000$ matrix
- Time

$$f * t_f + m * t_m = f * t_f * (1 + t_m/t_f * 1/q)$$

$$= 2 * n^2 * t_f * (1 + t_m/t_f * 1/2)$$

- For t_f and t_m , using data from R. Vuduc's PhD (pp 351-3)
 - <http://bebop.cs.berkeley.edu/pubs/vuduc2003-dissertation.pdf>
 - For t_m use minimum-memory-latency / words-per-cache-line

	Clock	Peak	Mem Lat (Min,Max)		Linesize	t _m /t _f
	MHz	Mflop/s	cycles		Bytes	
Ultra 2i	333	667	38	66	16	24.8
Ultra 3	900	1800	28	200	32	14.0
Pentium 3	500	500	25	60	32	6.3
Pentium3M	800	800	40	60	32	10.0
Power3	375	1500	35	139	128	8.8
Power4	1300	5200	60	10000	128	15.0
Itanium1	800	3200	36	85	32	36.0
Itanium2	900	3600	11	60	64	5.5


*machine
balance
(q must
be at least
this for
½ peak
speed)*

Simplifying Assumptions

- What simplifying assumptions did we make in this analysis?
 - Ignored parallelism in processor between memory and arithmetic within the processor
 - Sometimes drop arithmetic term in this type of analysis
 - Assumed fast memory was large enough to hold three vectors
 - Reasonable if we are talking about any level of cache
 - Not if we are talking about registers (~32 words)
 - Assumed the cost of a fast memory access is 0
 - Reasonable if we are talking about registers
 - Not necessarily if we are talking about cache (1-2 cycles for L1)
 - Memory latency is constant
- Could simplify even further by ignoring memory operations in X and Y vectors
 - Mflop rate/element = $2 / (2 * t_f + t_m)$


Validating the Model

- How well does the model predict actual performance?
 - Actual DGEMV: Most highly optimized code for the platform
- Model sufficient to compare across machines
- But under-predicting on most recent ones due to latency estimate



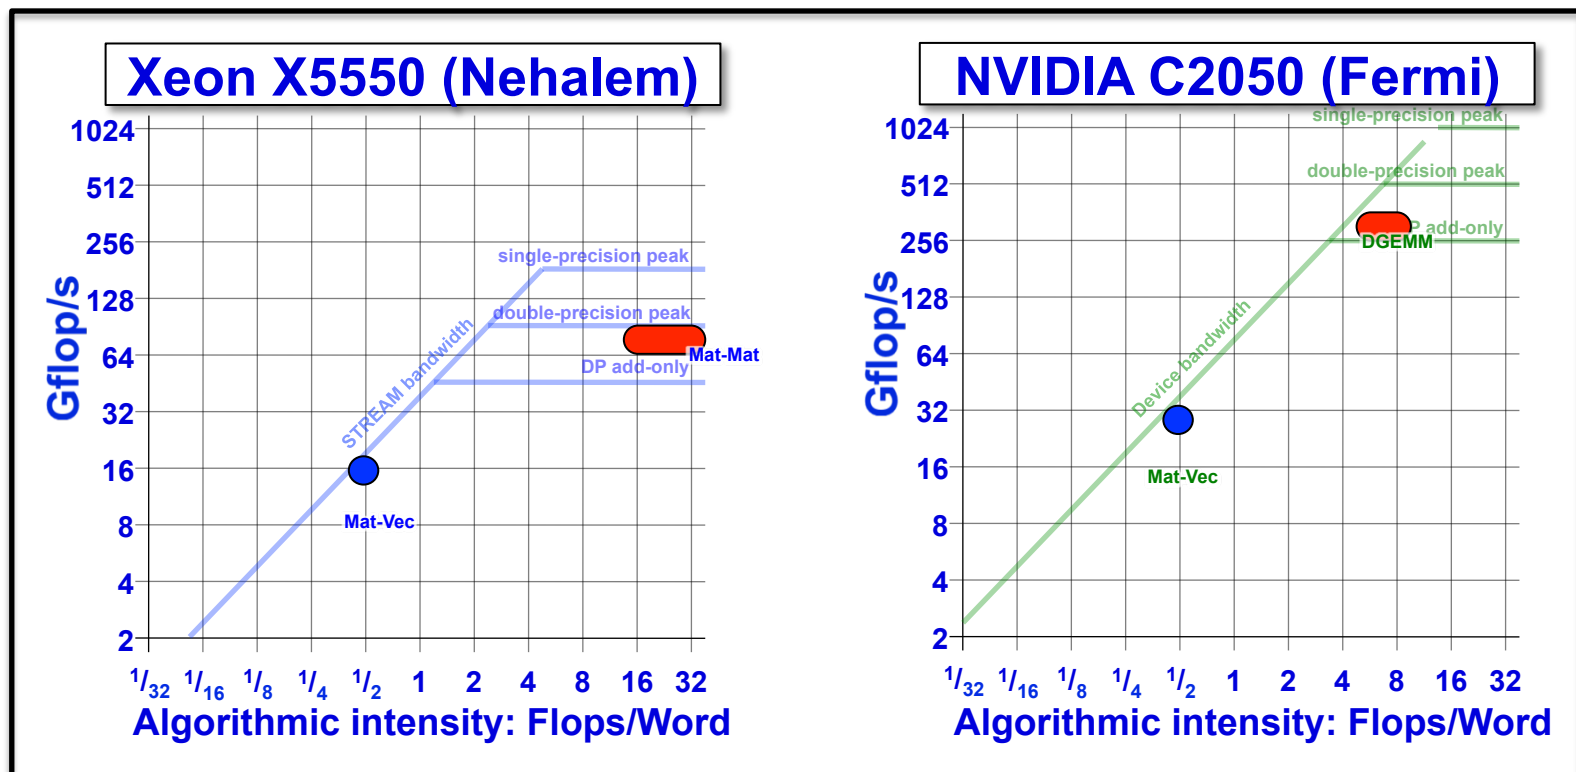
Roofline for Matrix-Vector multiply

- Roofline capture the gross differences between BLAS2 (matrix-vector) and BLAS3 (Matrix-Matrix)



Roofline for Matrix-Vector and Matrix-Matrix multiply

- Roofline capture the gross differences between BLAS2 (matrix-vector) and BLAS3 (Matrix-Matrix)



Naïve Matrix Multiply

```
{implements  $C = C + A*B$ }
```

```
for i = 1 to n
```


```
  for j = 1 to n
```

```
    for k = 1 to n
```

```
       $C(i,j) = C(i,j) + A(i,k) * B(k,j)$ 
```

Algorithm has $2*n^3 = O(n^3)$ Flops and
operates on $3*n^2$ words of memory

q potentially as large as $2*n^3 / 3*n^2 = O(n)$



Naïve Matrix Multiply

{implements $C = C + A*B$ }

for $i = 1$ to n

{read row i of A into fast memory}

for $j = 1$ to n


{read $C(i,j)$ into fast memory}

{read column j of B into fast memory}

for $k = 1$ to n

$C(i,j) = C(i,j) + A(i,k) * B(k,j)$

{write $C(i,j)$ back to slow memory}



Naïve Matrix Multiply


Number of slow memory references on unblocked matrix multiply

$$\begin{aligned} m &= n^3 && \text{to read each column of B } n \text{ times} \\ &+ n^2 && \text{to read each row of A once} \\ &+ 2n^2 && \text{to read and write each element of C once} \\ &= n^3 + 3n^2 \end{aligned}$$


$$\text{So } q = f / m = 2n^3 / (n^3 + 3n^2)$$

≈ 2 for large n , no improvement over matrix-vector multiply

Inner two loops are just matrix-vector multiply, of row i of A times B
Similar for any other order of 3 loops




Matrix-multiply, optimized several ways




Speed of n-by-n matrix multiply on Sun Ultra-1/170, peak = 330 MFlops

Naïve Matrix Multiply on RS/6000



$O(N^3)$ performance would have constant cycles/flop
Performance looks like $O(N^{4.7})$

Naïve Matrix Multiply on RS/6000



Blocked (Tiled) Matrix Multiply

Consider A,B,C to be N-by-N matrices of b-by-b subblocks where $b = n / N$ is called the **block size**

for i = 1 to N

for j = 1 to N

{read block C(i,j) into fast memory}

for k = 1 to N

{read block A(i,k) into fast memory}

{read block B(k,j) into fast memory}


C(i,j) = C(i,j) + A(i,k) * B(k,j) {do a matrix multiply on blocks}

{write block C(i,j) back to slow memory}

cache does this automatically

3 nested loops inside

block size = loop bounds



Tiling for registers (managed by you/compiler) or caches (hardware)

Blocked (Tiled) Matrix Multiply

Recall:

m is amount memory traffic between slow and fast memory

matrix has $n \times n$ elements, and $N \times N$ blocks each of size $b \times b$

f is number of floating point operations, $2n^3$ for this problem

$q = f / m$ is our measure of algorithm efficiency in the memory system

So:

$$\begin{aligned} m &= N * n^2 \quad \text{read each block of B } N^3 \text{ times } (N^3 * b^2 = N^3 * (n/N)^2 = N * n^2) \\ &+ N * n^2 \quad \text{read each block of A } N^3 \text{ times} \\ &+ 2n^2 \quad \text{read and write each block of C once} \\ &= (2N + 2) * n^2 \end{aligned}$$

$$\begin{aligned} \text{So computational intensity } q &= f / m = 2n^3 / ((2N + 2) * n^2) \\ &\approx n / N = b \quad \text{for large } n \end{aligned}$$

So we can improve performance by increasing the blocksize b

Can be much faster than matrix-vector multiply ($q=2$)

Using Analysis to Understand Machines

The blocked algorithm has computational intensity $q \approx b$

- The larger the block size, the more efficient our algorithm will be
- Limit: All three blocks from A,B,C must fit in fast memory (cache), so we cannot make these blocks arbitrarily large
- Assume your fast memory has size M_{fast}

$$3b^2 \leq M_{\text{fast}}, \quad \text{so} \quad q \approx b \leq (M_{\text{fast}}/3)^{1/2}$$

- To build a machine to run matrix multiply at 1/2 peak arithmetic speed of the machine, we need a fast memory of size

$$M_{\text{fast}} \geq 3b^2 \approx 3q^2 = 3(t_m/t_f)^2$$

- This size is reasonable for L1 cache, but not for register sets
- Note: analysis assumes it is possible to schedule the instructions perfectly

	t_m/t_f	required KB
Ultra 2i	24.8	14.8
Ultra 3	14	4.7
Pentium 3	6.25	0.9
Pentium3M	10	2.4
Power3	8.75	1.8
Power4	15	5.4
Itanium1	36	31.1
Itanium2	5.5	0.7

Limits to Optimizing Matrix Multiply

- The blocked algorithm changes the order in which values are accumulated into each $C[i,j]$ by applying commutativity and associativity
 - Get slightly different answers from naïve code, because of roundoff - OK
- The previous analysis showed that the blocked algorithm has computational intensity:

$$q \approx b \leq (M_{\text{fast}}/3)^{1/2}$$

- There is a lower bound result that says we cannot do any better than this (using only associativity, so still doing n^3 multiplications)
- **Theorem (Hong & Kung, 1981): Any reorganization of this algorithm (that uses only associativity) is limited to $q = O((M_{\text{fast}})^{1/2})$**
 - **#words moved between fast and slow memory = $\Omega(n^3 / (M_{\text{fast}})^{1/2})$**

Communication lower bounds for Matmul

- Hong/Kung theorem is a lower bound on amount of data communicated by matmul
 - Number of words moved between fast and slow memory (cache and DRAM, or DRAM and disk, or ...) = $\Omega(n^3 / M_{\text{fast}}^{1/2})$
- Cost of moving data may also depend on the number of “messages” into which data is packed
 - Eg: number of cache lines, disk accesses, ...
 - #messages = $\Omega(n^3 / M_{\text{fast}}^{3/2})$
- Lower bounds extend to anything “similar enough” to 3 nested loops
 - Rest of linear algebra (solving linear systems, least squares...)
 - Dense and sparse matrices
 - Sequential and parallel algorithms, ...
- More recent: extends to any nested loops accessing arrays
- Need (more) new algorithms to attain these lower bounds...

Summary of Lecture 2

- Details of machine are important for performance
 - Processor and memory system (not just parallelism)
 - Before you parallelize, make sure you're getting good serial performance
 - What to expect? Use understanding of hardware limits
- There is parallelism hidden within processors
 - Pipelining, SIMD, etc
- Machines have memory hierarchies
 - 100s of cycles to read from DRAM (main memory)
 - Caches are fast (small) memory that optimize average case
- Locality is at least as important as computation
 - Temporal: re-use of data recently used
 - Spatial: using data nearby to recently used data
- Can rearrange code/data to improve locality
 - Goal: minimize communication = data movement

Class Logistics

- Homework 0 posted on web site
 - Find and describe interesting application of parallelism
 - Due Wednesday Jan 24
 - Could even be your intended class project
 - We will add undergrads to bcourses so you can submit
- Fill in on-line class survey by midnight
 - We need this to assign teams for Homework 1
 - Teams will be announced no later than Monday
 - HW 1 will be posted by then