
CS 267


Lecture 9: Distributed Memory Machines and Programming

Aydin Buluc

<https://sites.google.com/lbl.gov/cs267-spr2018/>


- **Distributed Memory Architectures**
 - Properties of communication networks
 - Topologies
 - Performance models
- **Programming Distributed Memory Machines using Message Passing**
 - Overview of MPI
 - Basic send/receive use
 - Non-blocking communication
 - Collectives

Architectures in Top 500, Nov 2014



Historical Perspective

- **Early distributed memory machines were:**
 - Collection of microprocessors.
 - Communication was performed using bi-directional queues between nearest neighbors.
- **Messages were forwarded by processors on path.**
 - “Store and forward” networking
- **There was a strong emphasis on topology in algorithms, in order to minimize the number of hops = minimize time**



Network Analogy

To have a large number of different transfers occurring at once, you need a large number of distinct wires

- Not just a bus, as in shared memory

Networks are like streets:

- **Link** = street.
- **Switch** = intersection.
- **Distances** (hops) = number of blocks traveled.
- **Routing algorithm** = travel plan.

Properties:

Latency: how long to get between nodes in the network.

- Street: time for one car = $\text{dist (miles)} / \text{speed (miles/hr)}$

Bandwidth: how much data can be moved per unit time.

- Street: $\text{cars/hour} = \text{density (cars/mile)} * \text{speed (miles/hr)} * \text{\#lanes}$
- Network bandwidth is limited by the bit rate per wire and \#wires


Design Characteristics of a Network

- **Topology** (how things are connected)
 - Crossbar; ring; 2-D, 3-D, higher-D mesh or torus; hypercube; tree; butterfly; perfect shuffle, dragon fly, ...
- **Routing algorithm:**
 - Example in 2D torus: all east-west then all north-south (avoids deadlock).
- **Switching strategy:**
 - Circuit switching: full path reserved for entire message, like the telephone.
 - Packet switching: message broken into separately-routed packets, like the post office, or internet
- **Flow control** (what if there is congestion):
 - Stall, store data temporarily in buffers, re-route data to other nodes, tell source node to temporarily halt, discard, etc.

Performance Properties of a Network: Latency

- **Diameter:** the maximum (over all pairs of nodes) of the shortest path between a given pair of nodes.
- **Latency:** delay between send and receive times
 - Latency tends to vary widely across architectures
 - Vendors often report **hardware latencies** (wire time)
 - Application programmers care about **software latencies** (user program to user program)
- **Observations:**
 - Latencies differ by 1-2 orders across network designs
 - Software/hardware overhead at source/destination dominate cost (1s-10s usecs)
 - Hardware latency varies with distance (10s-100s nsec per hop) but is small compared to overheads
- **Latency is key for programs with many small messages**

End to End Latency (1/2 roundtrip) Over Time




- Latency has not improved significantly, unlike Moore's Law
 - T3E (shmem) was lowest point – in 1997

Data from Kathy Yelick


Performance Properties of a Network: Bandwidth

- The **bandwidth** of a link = # wires / time-per-bit
- Bandwidth typically in Gigabytes/sec (GB/s), i.e., 8×2^{20} bits per second
- **Effective bandwidth** is usually lower than physical link bandwidth due to packet overhead.

- Bandwidth is important for applications with mostly large messages



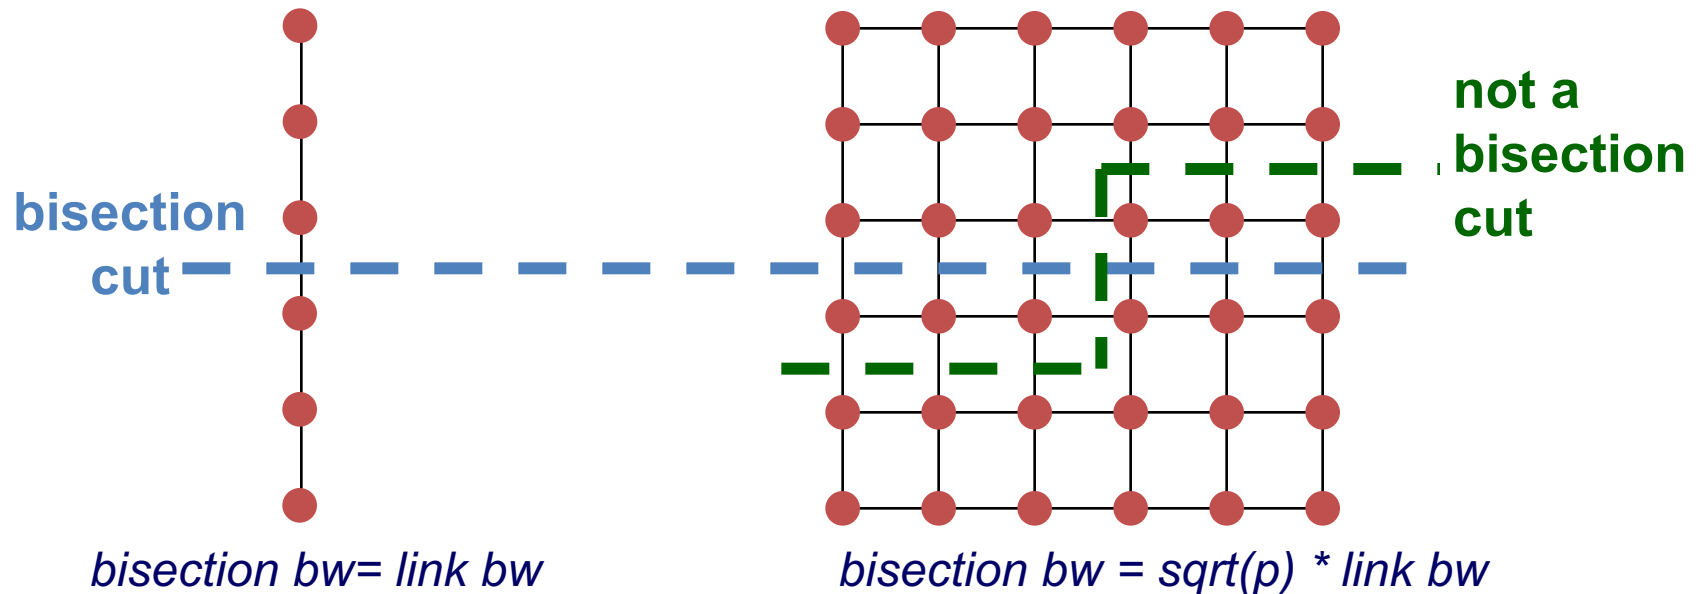
Bandwidth Chart



Data from Mike Welcome, NERSC

Performance Properties of a Network: Bisection Bandwidth

- **Bisection bandwidth:** bandwidth across smallest cut that divides network into two equal halves
- Bandwidth across “narrowest” part of the network



- Bisection bandwidth is important for algorithms in which all processors need to communicate with all others

Network Topology

- In the past, there was considerable research in network topology and in mapping algorithms to topology.
 - Key cost to be minimized: number of “hops” between nodes (e.g. “store and forward”)
 - Modern networks hide hop cost (i.e., “wormhole routing”), so topology less of a factor in performance of many algorithms
- Example: On IBM SP system, hardware latency varies from 0.5 usec to 1.5 usec, but user-level message passing latency is roughly 36 usec.
- Need some background in network topology
 - Algorithms may have a communication topology
 - Example later of big performance impact



Linear and Ring Topologies

◦ Linear array



- Diameter = $n-1$; average distance $\sim n/3$.
- Bisection bandwidth = 1 (in units of link bandwidth).

◦ Torus or Ring




- Diameter = $n/2$; average distance $\sim n/4$.
- Bisection bandwidth = 2.
- Natural for algorithms that work with 1D arrays.

Meshes and Tori – used in Hopper


Two dimensional mesh

- Diameter = $2 * (\text{sqrt}(n) - 1)$
- Bisection bandwidth = $\text{sqrt}(n)$



Two dimensional torus

- Diameter = $\text{sqrt}(n)$
- Bisection bandwidth = $2 * \text{sqrt}(n)$




- Generalizes to higher dimensions
 - Cray XT (eg Hopper@NERSC) uses 3D Torus
- Natural for algorithms that work with 2D and/or 3D arrays (matmul)

Hypercubes

◦ **Number of nodes $n = 2^d$ for dimension d .**

- Diameter = d .
- Bisection bandwidth = $n/2$.




◦ **Popular in early machines (Intel iPSC, NCUBE).**

- Lots of clever algorithms.
- See 1996 online CS267 notes.




◦ **Greycode addressing:**

- Each node connected to d others with 1 bit different.



Trees


- **Diameter = $\log n$.**
- **Bisection bandwidth = 1.**
- **Easy layout as planar graph.**
- **Many tree algorithms (e.g., summation).**
- **Fat trees avoid bisection bandwidth problem:**
 - More (or wider) links near top.
 - Example: Thinking Machines CM-5.




Butterflies

- Diameter = $\log n$.
- Bisection bandwidth = n .
- Cost: lots of wires.
- Used in BBN Butterfly.
- Natural for FFT.

Ex: to get from proc 101 to 110,
Compare bit-by-bit and
Switch if they disagree, else not




butterfly switch



multistage butterfly network

Does Topology Matter?



See EECS Tech Report *UCB/EECS-2011-92*, August 2011

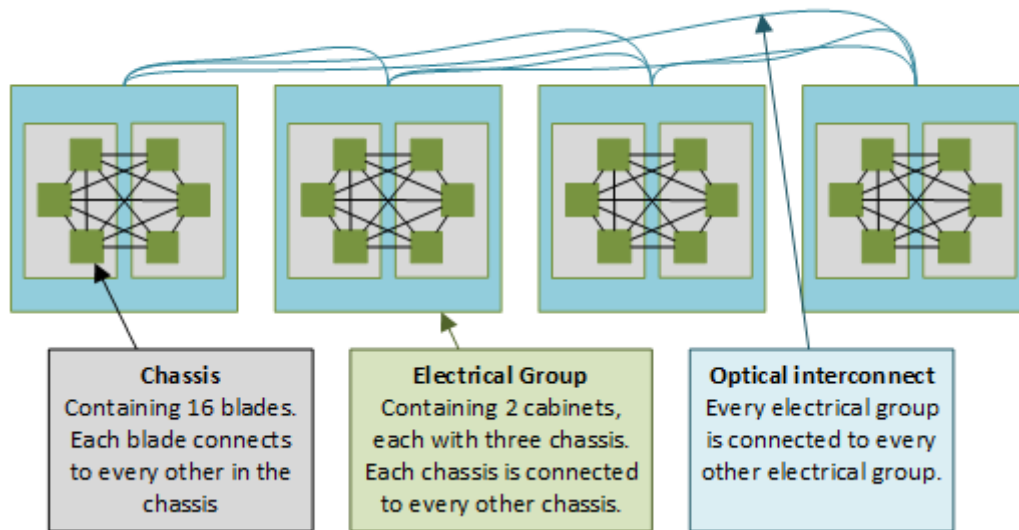
Dragonflies – used in Edison and Cori

- **Motivation:** Exploit gap in cost and performance between optical interconnects (which go between cabinets in a machine room) and electrical networks (inside cabinet)
 - Optical more expensive but higher bandwidth when long
 - Electrical networks cheaper, faster when short
- **Combine in hierarchy:**
 - Several groups are connected together using all to all links, i.e. each group has at least one link directly to each other group.
 - The topology inside each group can be any topology.
- **Uses a randomized routing algorithm**
- **Outcome:** programmer can (usually) ignore topology, get good performance
 - Important in virtualized, dynamic environment
 - Drawback: variable performance

“Technology-Drive, Highly-Scalable Dragonfly Topology,” ISCA 2008


Dragonfly in practice

Aries interconnect for an 8 cabinet Cray XC40



Source: European Centre for Medium-Range Weather Forecasts

Source of image on the right (and more info):
<http://www.nersc.gov/users/computational-systems/edison/configuration/interconnect/>



Performance Models

Shared Memory Performance Models

- **Parallel Random Access Memory (PRAM)**
- **All memory access operations complete in one clock period -- no concept of memory hierarchy (“too good to be true”).**
 - OK for understanding whether an algorithm has enough parallelism at all (see CS273).
 - Parallel algorithm design strategy: first do a PRAM algorithm, then worry about memory/communication time (sometimes works)
- **Slightly more realistic versions exist**
 - E.g., Concurrent Read Exclusive Write (CREW) PRAM.
 - Still missing the memory hierarchy

Latency and Bandwidth Model

- Time to send message of length n is roughly

$$\begin{aligned}\text{Time} &= \text{latency} + n * \text{cost_per_word} \\ &= \text{latency} + n / \text{bandwidth}\end{aligned}$$

- Topology is assumed irrelevant.
- Often called “ α – β model” and written

$$\text{Time} = \alpha + n * \beta$$


- Usually $\alpha \gg \beta \gg \text{time per flop}$.
 - One long message is cheaper than many short ones.
 - Can do hundreds or thousands of flops for cost of one message.

$$\alpha + n * \beta \ll n * (\alpha + 1 * \beta)$$



- Lesson: Need large computation-to-communication ratio to be efficient.
- LogP – more detailed model (Latency/overhead/gap/Proc.)

Latency and Bandwidth Model in 2010 and beyond+

Processors are multi-core and many nodes are multi-chip



NIC bandwidth bottleneck



- Ping-pong between 2 nodes using 1-16 cores on each node
 - Top: BG/Q,
 - Bottom Cray XE6
 - X-axis is message size
-
- Alpha-beta model predicts a single curve – rates independent of the number of communicating processes

Programming Distributed Memory Machines with Message Passing

Slides by

Aydin Buluc, Jonathan Carter, Jim Demmel,
Bill Gropp, Kathy Yelick

Message Passing Libraries (1)

- Many “message passing libraries” were once available
 - Chameleon, from ANL.
 - CMMD, from Thinking Machines.
 - Express, commercial.
 - MPL, native library on IBM SP-2.
 - NX, native library on Intel Paragon.
 - Zipcode, from LLL.
 - PVM, Parallel Virtual Machine, public, from ORNL/UTK.
 - Others...
 - **MPI, Message Passing Interface, now the industry standard.**
- Need standards to write portable code.

Message Passing Libraries (2)

- **All communication, synchronization require subroutine calls**
 - **No shared variables**
 - **Program run on a single processor just like any uniprocessor program, except for calls to message passing library**
- **Subroutines for**
 - **Communication**
 - **Pairwise or point-to-point: Send and Receive**
 - **Collectives all processor get together to**
 - Move data: Broadcast, Scatter/gather
 - Compute and move: sum, product, max, prefix sum, ... of data on many processors
 - **Synchronization**
 - **Barrier**
 - **No locks because there are no shared variables to protect**
 - **Enquiries**
 - **How many processes? Which one am I? Any messages waiting?**

Novel Features of MPI




- Communicators encapsulate communication spaces for library safety
- Datatypes reduce copying costs and permit heterogeneity
- Multiple communication modes allow precise buffer management
- Extensive collective operations for scalable global communication
- Process topologies permit efficient process placement, user views of process layout
- Profiling interface encourages portable tools

MPI References

- **The Standard itself:**
 - at <http://www.mpi-forum.org>
 - All MPI official releases, in both postscript and HTML
 - Latest version MPI 3.1, released June 2015
- **Other information on Web:**
 - at <http://www.mcs.anl.gov/research/projects/mpi/index.htm>
 - pointers to lots of stuff, including other talks and tutorials, a FAQ, other MPI pages

Books on MPI

- ***Using MPI: Portable Parallel Programming with the Message-Passing Interface (third edition)***, by Gropp, Lusk, and Skjellum, MIT Press, 2014.
- ***Using Advanced MPI: Modern Features of the Message-Passing Interface***, by Gropp, Hoefler, Thakur, and Lusk, MIT Press, 2014
- ***Using MPI-2: Portable Parallel Programming with the Message-Passing Interface***, by Gropp, Lusk, and Thakur, MIT Press, 1999.
- ***MPI: The Complete Reference - Vol 1 The MPI Core***, by Snir, Otto, Huss-Lederman, Walker, and Dongarra, MIT Press, 1998.
- ***MPI: The Complete Reference - Vol 2 The MPI Extensions***, by Gropp, Huss-Lederman, Lumsdaine, Lusk, Nitzberg, Saphir, and Snir, MIT Press, 1998.
- ***Designing and Building Parallel Programs***, by Ian Foster, Addison-Wesley, 1995.
- ***Parallel Programming with MPI***, by Peter Pacheco, Morgan-Kaufmann, 1997.



Finding Out About the Environment

- Two important questions that arise early in a parallel program are:
 - How many processes are participating in this computation?
 - Which one am I?
- MPI provides functions to answer these questions:
 - `MPI_Comm_size` reports the number of processes.
 - `MPI_Comm_rank` reports the *rank*, a number between 0 and size-1, identifying the calling process

Hello (C)

```
#include "mpi.h"
#include <stdio.h>


int main( int argc, char *argv[] )
{
    int rank, size;
    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    MPI_Comm_size( MPI_COMM_WORLD, &size );
    printf( "I am %d of %d\n", rank, size );
    MPI_Finalize();
    return 0;
}
```

Notes on Hello World

- All MPI programs begin with `MPI_Init` and end with `MPI_Finalize`
- `MPI_COMM_WORLD` is defined by `mpi.h` (in C) or `mpif.h` (in Fortran) and designates all processes in the MPI “job”
- Each statement executes independently in each process
 - including the `printf/print` statements
- The MPI-1 Standard does not specify how to run an MPI program, but many implementations provide `mpirun -np 4 a.out`

MPI Basic Send/Receive

- **We need to fill in the details in**



- **Things that need specifying:**
 - How will “data” be described?
 - How will processes be identified?
 - How will the receiver recognize/screen messages?
 - What will it mean for these operations to complete?

Some Basic Concepts

- ° Processes can be collected into groups
- ° Each message is sent in a context, and must be received in the same context
 - Provides necessary support for libraries
- ° A group and context together form a communicator
- ° A process is identified by its rank in the group associated with a communicator
- ° There is a default communicator whose group contains all initial processes, called `MPI_COMM_WORLD`


MPI Datatypes

- The data in a message to send or receive is described by a triple (address, count, datatype), where
- An MPI datatype is recursively defined as:
 - predefined, corresponding to a data type from the language (e.g., MPI_INT, MPI_DOUBLE)
 - a contiguous array of MPI datatypes
 - a strided block of datatypes
 - an indexed array of blocks of datatypes
 - an arbitrary structure of datatypes
- There are MPI functions to construct custom datatypes, in particular ones for subarrays
- May hurt performance if datatypes are complex

MPI Tags

- **Messages are sent with an accompanying user-defined integer tag, to assist the receiving process in identifying the message**
- **Messages can be screened at the receiving end by specifying a specific tag, or not screened by specifying `MPI_ANY_TAG` as the tag in a receive**
- **Some non-MPI message-passing systems have called tags “message types”. MPI calls them tags to avoid confusion with datatypes**

MPI Basic (Blocking) Send




`MPI_Send(A, 10, MPI_DOUBLE, 1, ...)`

`MPI_Recv(B, 20, MPI_DOUBLE, 0, ...)`

`MPI_SEND(start, count, datatype, dest, tag, comm)`

- The message buffer is described by (start, count, datatype).
- The target process is specified by dest, which is the rank of the target process in the communicator specified by comm.
- When this function returns, the data has been delivered to the system and the buffer can be reused. The message may not have been received by the target process.

MPI Basic (Blocking) Receive



`MPI_Send(A, 10, MPI_DOUBLE, 1, ...)`

`MPI_Recv(B, 20, MPI_DOUBLE, 0, ...)`

`MPI_RECV(start, count, datatype, source, tag, comm, status)`

- **Waits until a matching (both source and tag) message is received from the system, and the buffer can be used**
- **source** is rank in communicator specified by **comm**, or **MPI_ANY_SOURCE**
- **tag** is a tag to be matched or **MPI_ANY_TAG**
- **receiving fewer than count occurrences of datatype is OK, but receiving more is an error**
- **status** contains further information (e.g. size of message)

A Simple MPI Program

```
#include "mpi.h"
#include <stdio.h>
int main( int argc, char *argv[])
{
    int rank, buf;
    MPI_Status status;
    MPI_Init(&argv, &argc);
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );

    /* Process 0 sends and Process 1 receives */
    if (rank == 0) {
        buf = 123456;
        MPI_Send( &buf, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
    }
    else if (rank == 1) {
        MPI_Recv( &buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD,
                  &status );
        printf( "Received %d\n", buf );
    }

    MPI_Finalize();
    return 0;
}
```

Retrieving Further Information


- **Status** is a data structure allocated in the user's program.
- **In C:**

```
int recvd_tag, recvd_from, recvd_count;
MPI_Status status;
MPI_Recv(..., MPI_ANY_SOURCE, MPI_ANY_TAG, ..., &status )
recvd_tag  = status.MPI_TAG;
recvd_from = status.MPI_SOURCE;
MPI_Get_count( &status, datatype, &recvd_count );
```

MPI can be simple

- **Claim: most MPI applications can be written with only 6 functions (although which 6 may differ)**
- Using point-to-point:
 - `MPI_INIT`
 - `MPI_FINALIZE`
 - `MPI_COMM_SIZE`
 - `MPI_COMM_RANK`
 - `MPI_SEND`
 - `MPI_RECEIVE`
- Using collectives:
 - `MPI_INIT`
 - `MPI_FINALIZE`
 - `MPI_COMM_SIZE`
 - `MPI_COMM_RANK`
 - `MPI_BCAST`
 - `MPI_REDUCE`
- **You may use more for convenience or performance**

PI redux: Numerical integration



Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$


We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width Δx and height $F(x_i)$ at the middle of interval i .

Example: Calculating Pi

E.g., in a 4-process run, each process gets every 4th interval. Process 0 slices are in red.



- **Simple program written in a data parallel style in MPI**
 - E.g., for a reduction (recall “tricks with trees” lecture), each process will first reduce (sum) its own values, then call a collective to combine them
- **Estimates pi by approximating the area of the quadrant of a unit circle**
- **Each process gets $1/p$ of the intervals (mapped round robin, i.e., a cyclic mapping)**

Example: PI in C – 1/2

```
#include "mpi.h"
#include <math.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    int done = 0, n, myid, numprocs, i, rc;
    double PI25DT = 3.141592653589793238462643;
    double mypi, pi, h, sum, x, a;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    while (!done) {
        if (myid == 0) {
            printf("Enter the number of intervals: (0 quits) ");
            scanf("%d", &n);
        }
        MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
        if (n == 0) break;
    }
}
```

Example: PI in C – 2/2

```
h = 1.0 / (double) n;
sum = 0.0;
for (i = myid + 1; i <= n; i += numprocs) {
    x = h * ((double)i - 0.5);
    sum += 4.0 * sqrt(1.0 - x*x);
}
mypi = h * sum;
MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0,
           MPI_COMM_WORLD);
if (myid == 0)
    printf("pi is approximately %.16f, Error is .16f\n",
           pi, fabs(pi - PI25DT));
}
MPI_Finalize();

return 0;

}
```